

Diagnostic Tooling – Be ready when things go wrong in production



Aka, why you should be interested in the Node.js Diagnostics WG

NodeConfEU 2018 – Michael Dawson and Mike Kaufman

Loves the web and building software (with Node.js!)



Senior Software Developer @ IBM

IBM Node.js Community Lead

Node.js TSC chair and CommComm member

Active in many Node.js project teams, working groups and initiatives

Contact me:

michael dawson@ca.ibm.com

Github: @mhdawson

Twitter: [@mhdawson1](#)

<https://www.linkedin.com/in/michael-dawson-6051282>



About **Mike Kaufman**

Hates long plane rides, has a 4-year old

Principal Software Engineer @ Microsoft

Currently:

- Chakra JavaScript Engine
- Edge Browser
- Node.js Diagnostics Working Group

Previously:

- SQL Server, SQL Azure DB
- Visual Studio
- BEA Systems/WebLogic

Contact me:

Mike.Kaufman@microsoft.com



Agenda

- Fail Fast?
- Intro to Diagnostic Working Group
- Key WG Initiatives
- Scenarios/Demos

Fail Fast?

- Fail Fast and Restart?

<https://www.martinfowler.com/ieeeSoftware/failFast.pdf>

- Reasons to investigate
 - Accumulation
 - Early warning
 - Customer Impacting
 - Performance
 - Resource Usage (\$\$\$)
- Key metrics
 - Mean time to Failure
 - Mean time to Discovery
 - Mean time to Recovery

Scenarios

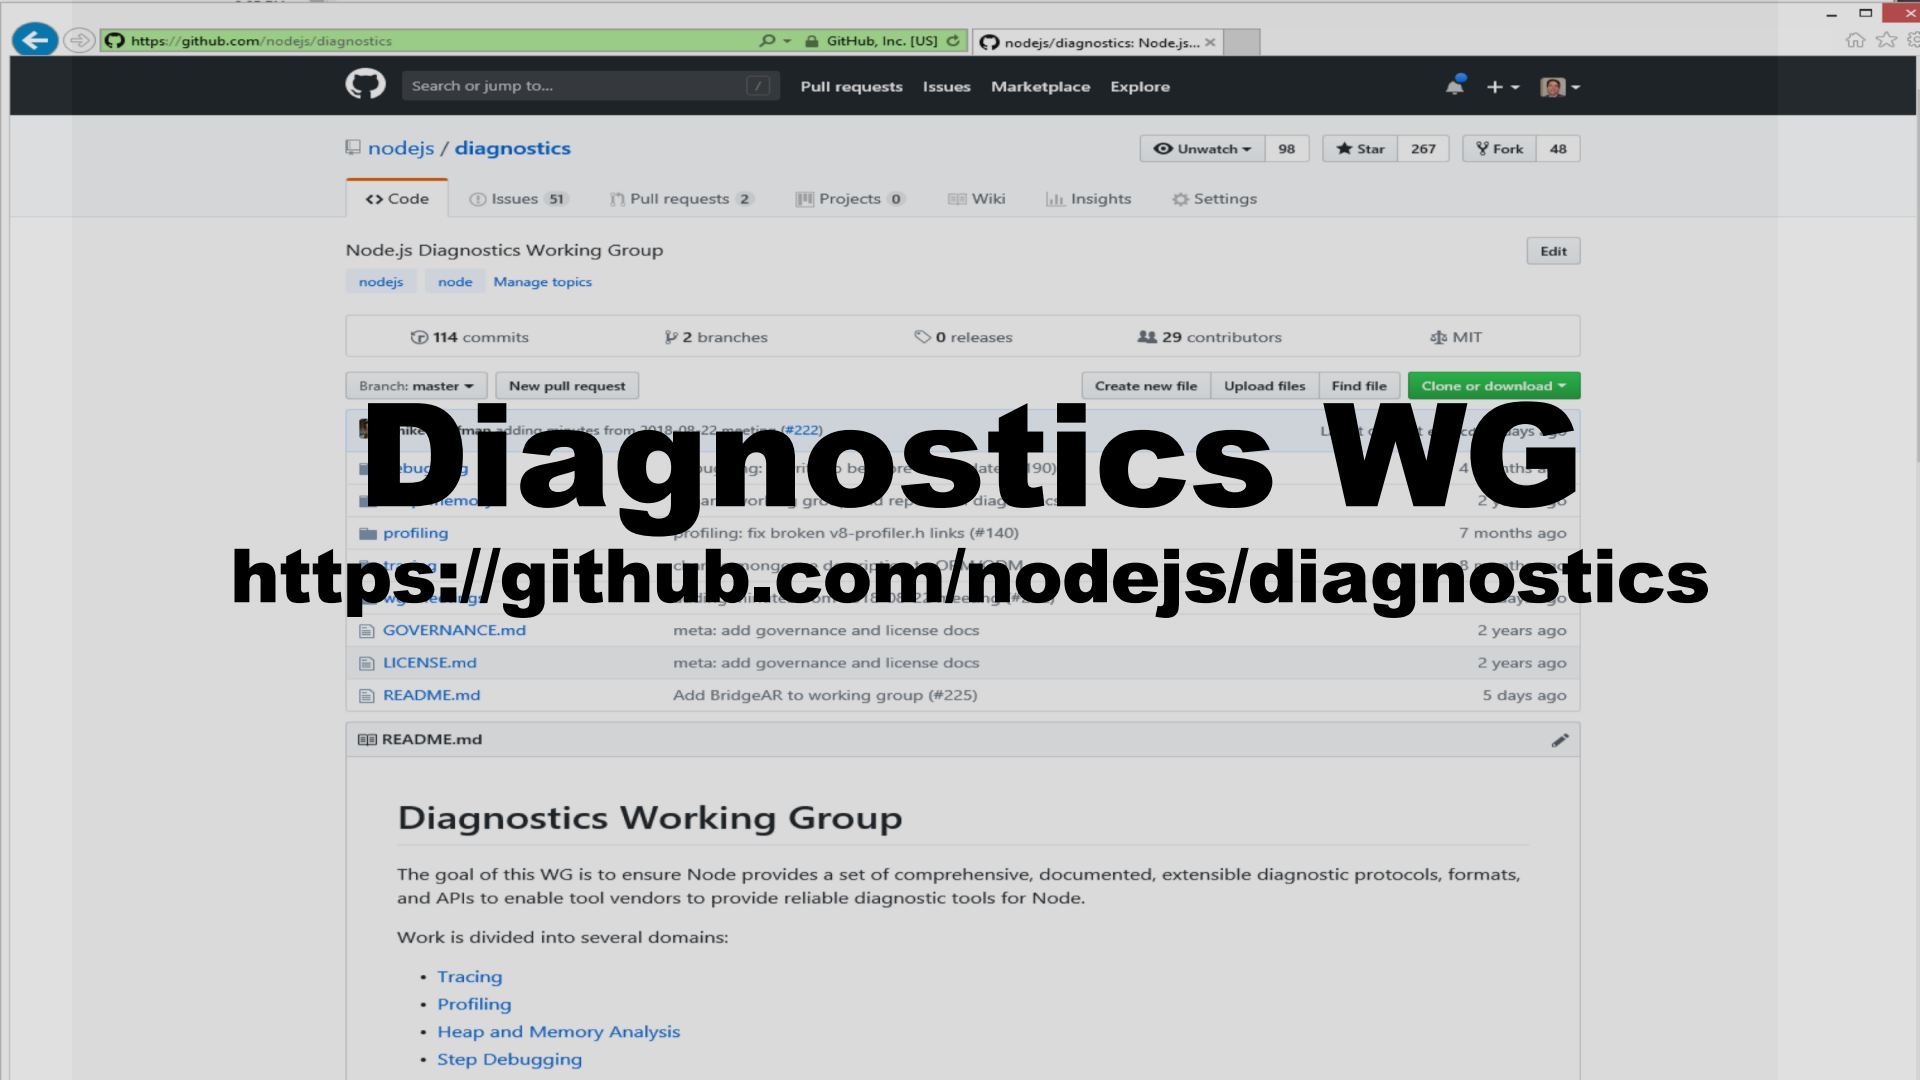
- Failure
 - Memory Leak
 - Crash
- Performance
 - Slow consistently/periodically
 - Does not scale
- Excessive resource usage
 - CPU
 - Memory
- Monitoring
 - Operational Data Collection
 - Transaction Data Collection

Toolkit

- Logging/Tracing
- Profilers
- Snapshot of execution/system state
 - First Failure Data Capture (FFDC)
 - Heap Dump
 - Core Dump
- APIs

Diagnostics WG

<https://github.com/nodejs/diagnostics>





Diagnostics WG mission

The goal of this WG is to ensure Node provides a set of comprehensive, documented, extensible diagnostic protocols, formats, and APIs to enable tool vendors to provide reliable diagnostic tools for Node.



Team Members

- [@AndreasMadsen](#) - Andreas Madsen
- [@BridgeAR](#) - Ruben Bridgewater
- [@cjihrig](#) - Colin Ihrig
- [@danielkhan](#) - Daniel Khan
- [@Flarna](#) - Gerhard Stöbich
- [@hashseed](#) - Yang Guo
- [@hekike](#) - Peter Marton
- [@Hollerberg](#) - Gernot Reisinger
- [@jasnell](#) - James M Snell
- [@jkrems](#) - Jan Olaf Krems
- [@joyeecheung](#) - Joyee Cheung
- [@kjin](#) - Kelvin Jin

- [@mcollina](#) - Matteo Collina
- [@mhdawson](#) - Michael Dawson
- [@mike-kaufman](#) - Mike Kaufman
- [@mmarchini](#) - Matheus Marchini
- [@naugtur](#) - Zbyszek Tenerowicz
- [@ofrobots](#) - Ali Ijaz Sheikh
- [@othiym23](#) - Forrest L Norvell
- [@Qard](#) - Stephen Belanger
- [@sam-github](#) - Sam Roberts
- [@watson](#) - Thomas Watson
- [@yunong](#) - Yunong Xiao



Key Initiatives

- Documentation/Best Practices
- Tooling/Data Sources
 - Support Tiers
 - Trace Events
 - node-report
 - Post mortem with Promises
 - Innode/JS API (@joyeecheung, @mmarchini)
 - Diagnostic Channel
- APIs
 - Async Hooks

Strategic Initiatives

Initiative	Champion	Stakeholders
Diagnostic Channel	@qard	
Async Hooks	@ofrobots	
Async Context	@mike-kaufman	@kjin
Node-report in core	@mhdawson	@richardlau
Support tiers	@mhdawson	
CPU Profiling	@mmarchini	
Post-mortem WG merge	@mmarchini	



Doc/Best Practices

- Diagnostics of no use UNLESS
 - People know they exist
 - People understand how to use them
- Some existing WG work (@naughtur)
 - <https://github.com/nodejs/diagnostics/issues/211>
 - <https://github.com/naughtur/node-diagnostics-howtos>
 - <https://github.com/nodejs/nodejs.org/pull/1444>



Tooling – Support Tiers

- Tiers
 - 1,2,3,4
 - Uncategorized
- Need help adding tests

Tool Type	Tool/API Name	Regular Testing in Node.js CI	Integrated with Node.js	Target Tier
FFDC	node-report	No	No	1
Memory	mdb_V8	No	No	4
Memory	node-heapdump	No	No	2
Memory	V8 heap profiler	No	Yes	1
Memory	V8 sampling heap profiler	No	Yes	1
AsyncFlow	Async Hooks (API)	?	Yes	1
Debugger	V8 Debug protocol (API)	No	Yes	1
Debugger	Command line Debug Client	?	Yes	1
Debugger	llnode	?	No	2
Debugger	Chrome Dev tools	?	No	3
Debugger	Chakracore - time-travel	No	Data source only	too early
Tracing	trace_events (API)	No	Yes	1
Tracing	DTrace	No	Partial	3
Tracing	LTTng	No	Removed?	N/A

<https://github.com/nodejs/node/blob/master/doc/guides/diagnostic-tooling-support-tiers.md>

APIs



- Async Context

- Formalization - <https://github.com/nodejs/diagnostics/tree/master/async-context/slide-decks>

- Async Hooks

```
function init(asyncId, type, triggerAsyncId, resource) { }  
function before(asyncId) { }  
function after(asyncId) { }  
function destroy(asyncId) { }  
function promiseResolve(asyncId) { }
```

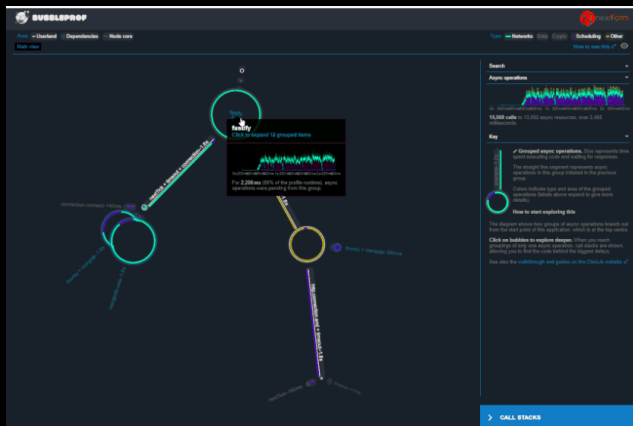
```
const asyncHook = async_hooks.createHook({ init, before, after, destroy, promiseResolve });
```

```
asyncHook.enable();  
asyncHook.disable();
```

FSEVENTWRAP, FSREQWRAP, GETADDRINFOREQWRAP, GETNAMEINFOREQWRAP, HTTPPARSER,
JSSTREAM, PIPECONNECTWRAP, PIPEWRAP, PROCESSWRAP, QUERYWRAP, SHUTDOWNWRAP,
SIGNALWRAP, STATWATCHER, TCPCONNECTWRAP, TCPSERVER, TCPWRAP, TIMERWRAP, TTYWRAP,
UDPSENDWRAP, UDPWRAP, WRITWRAP, ZLIB, SSLCONNECTION, PBKDF2REQUEST,
RANDOMBYTESREQUEST, TLSWRAP, Timeout, Immediate, TickObject

Some Examples

- BubbleProf



- Doctrine – HTTP Association
 - <https://blog.doctrine.fr/how-doctrine-has-massively-improved-its-node-js-error-analysis-with-async-hooks-2917e4b69d5d?gi=d0820a193d6d>

Ok, but how does this help ME?

Back to the Scenarios....

- Failure
 - Memory Leak
 - Crash
- Performance
 - Slow consistently/periodically
 - Does not scale
- Excessive resource usage
 - CPU
 - Memory
- Monitoring
 - Operation Data Collection
 - Transaction Data Collection

Scenario 1 – Node Report

- npm install node-report
- Easy, human readable dump
- node -r node-report test.js

node-report

2.2.1 • Public

```
export NODEREPORT_EVENTS=exception+fatalerror+signal+apicall  
export NODEREPORT_SIGNAL=SIGUSR2|SIGQUIT  
export NODEREPORT_FILENAME=stdout|stderr|<filename>  
export NODEREPORT_DIRECTORY=<full path>  
export NODEREPORT_VERBOSE=yes|no
```

<https://github.com/nodejs/node-report>

<https://developer.ibm.com/node/2016/08/18/nodereport-first-failure-data-capture-for-node-js/>

<https://developer.ibm.com/node/2017/02/14/announcing-node-report-version-2-1-0-essential-diagnostic-node-js-applications/>

<https://github.com/RuntimeTools/appmetrics-dash>

Scenario 1 – Node Report Not Enough Memory

```
function useMemory() {  
  const objectRoot = new Object();  
  for (i = 0; i < 1000000000; i++) {  
    objectRoot[i.toString()] = new Object();  
  }  
}  
  
useMemory();
```

```
export NODEREPORT_EVENTS=exception+fatalerror+signal+apicall  
node --max-old-space-size=128 -r node-report test.js
```

Demo

Scenario 2 – What is my app doing

- Basic tracing

- trace
 - trace-gc
 - trace-opt

- Trace events

- Trace engine from V8
 - Unified view across V8/Node
 - chrome://tracing
 - <https://nodejs.org/api/tracing.html>

<https://vimeo.com/287708597> - Gireesh - Follow your code: Node/V8 Tracing

Scenario 2 – What is my app doing

node --trace-events-categories v8,node.fs.sync test-file.js

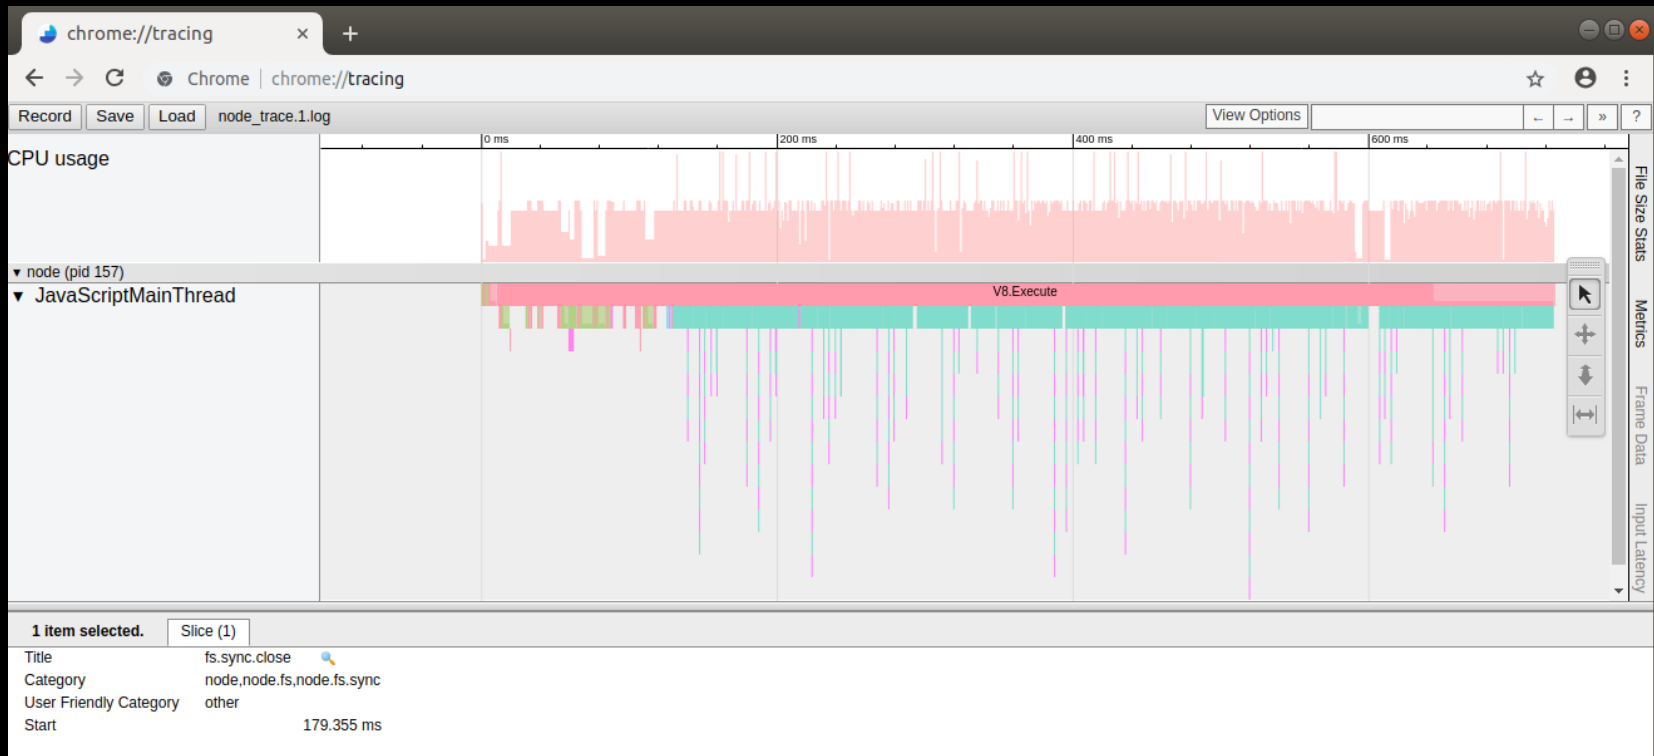
```
const fs = require('fs');
for (i = 0; i < 10000; i++) {
  const fd = fs.openSync('TestFile' + i, 'a');
  fs.closeSync(fd);
}
```

Scenario 2 – What is my app doing

node_trace.1.log

[illegible]

Scenario 2 – What is my app doing



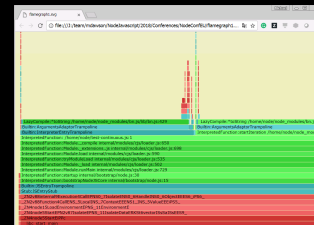
Scenario 3 – Performance Flame Graphs

- Where is my app spending CPU cycles?
- Will be part of best practices, some current info

<https://github.com/nodejs/nodejs.org/pull/1444>

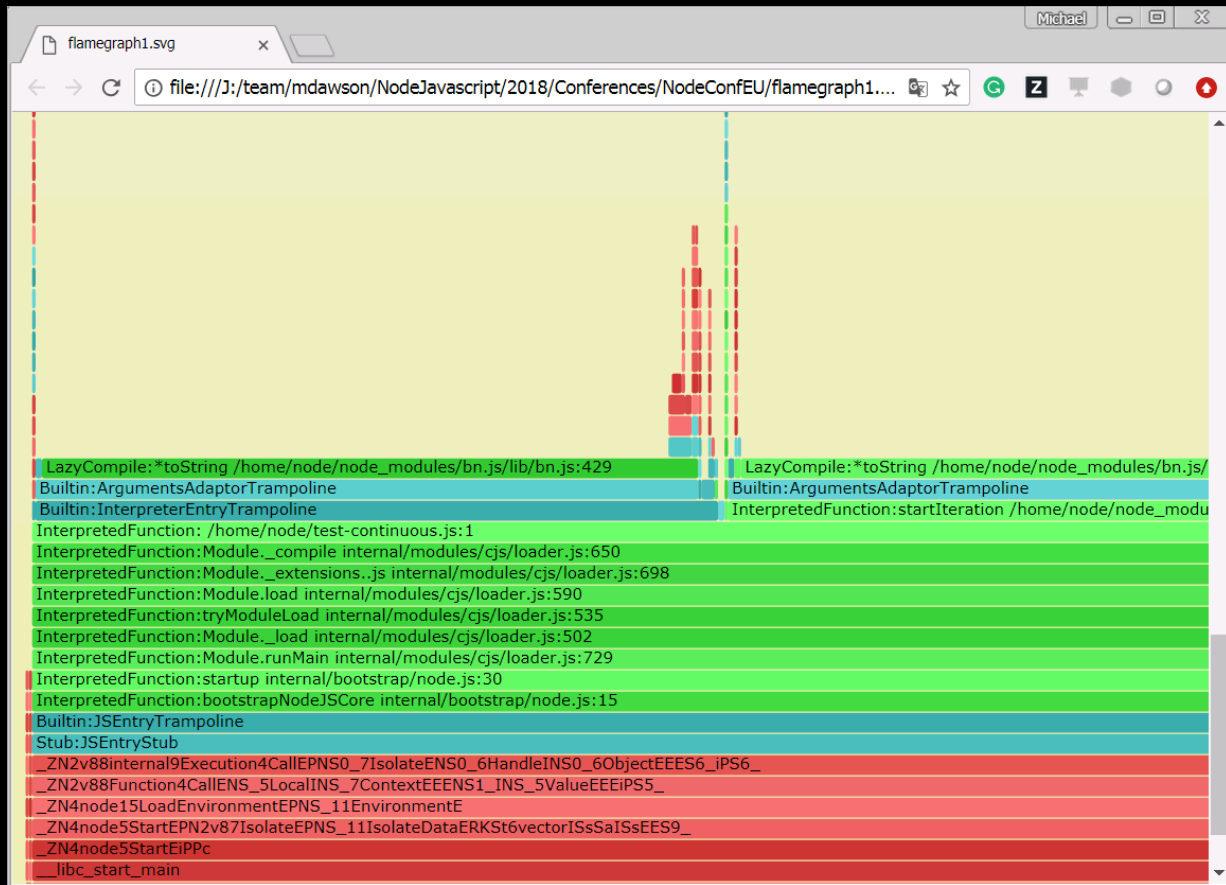
<https://github.com/mmmarchini/nodejs-production-diagnostic-tools#linux-perf>

- Will demo perf, also V8 profiler and tools that visualize.

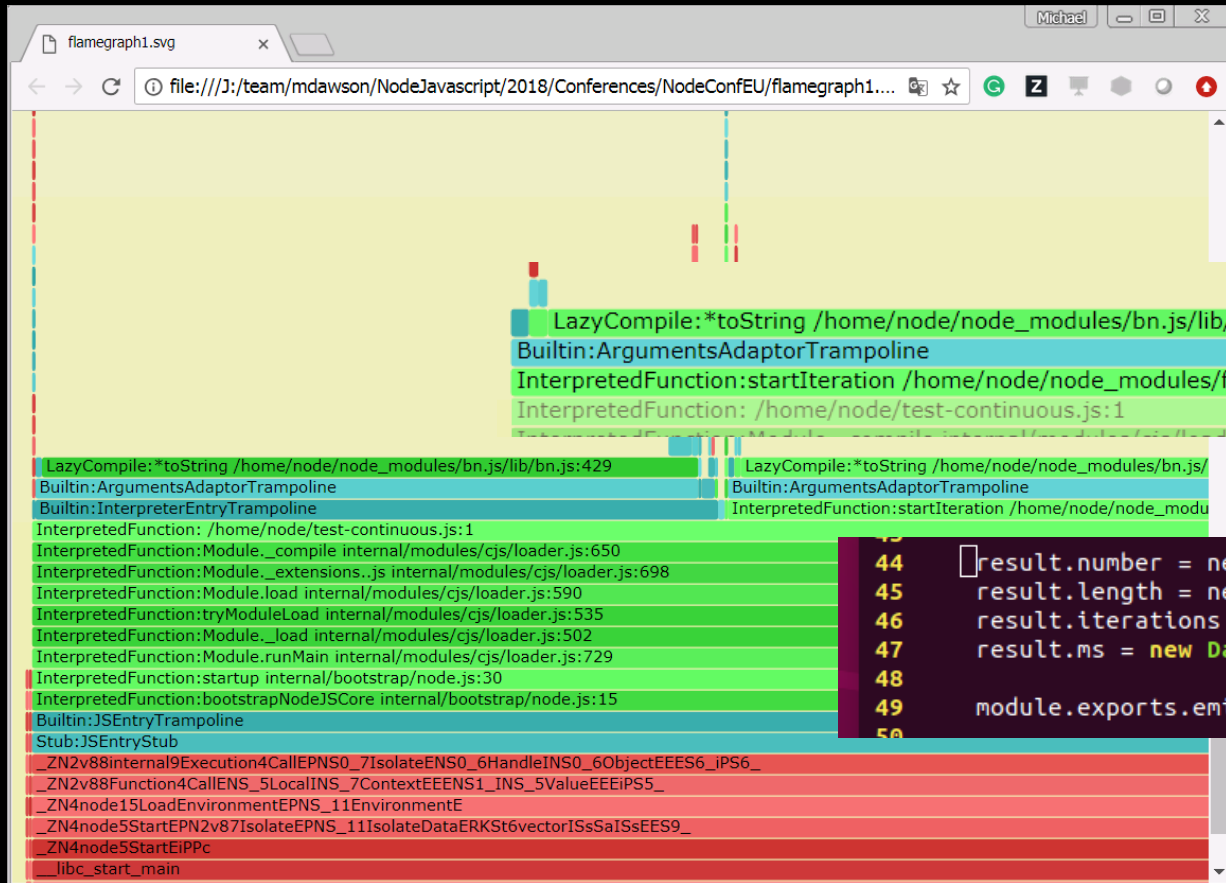


Demo

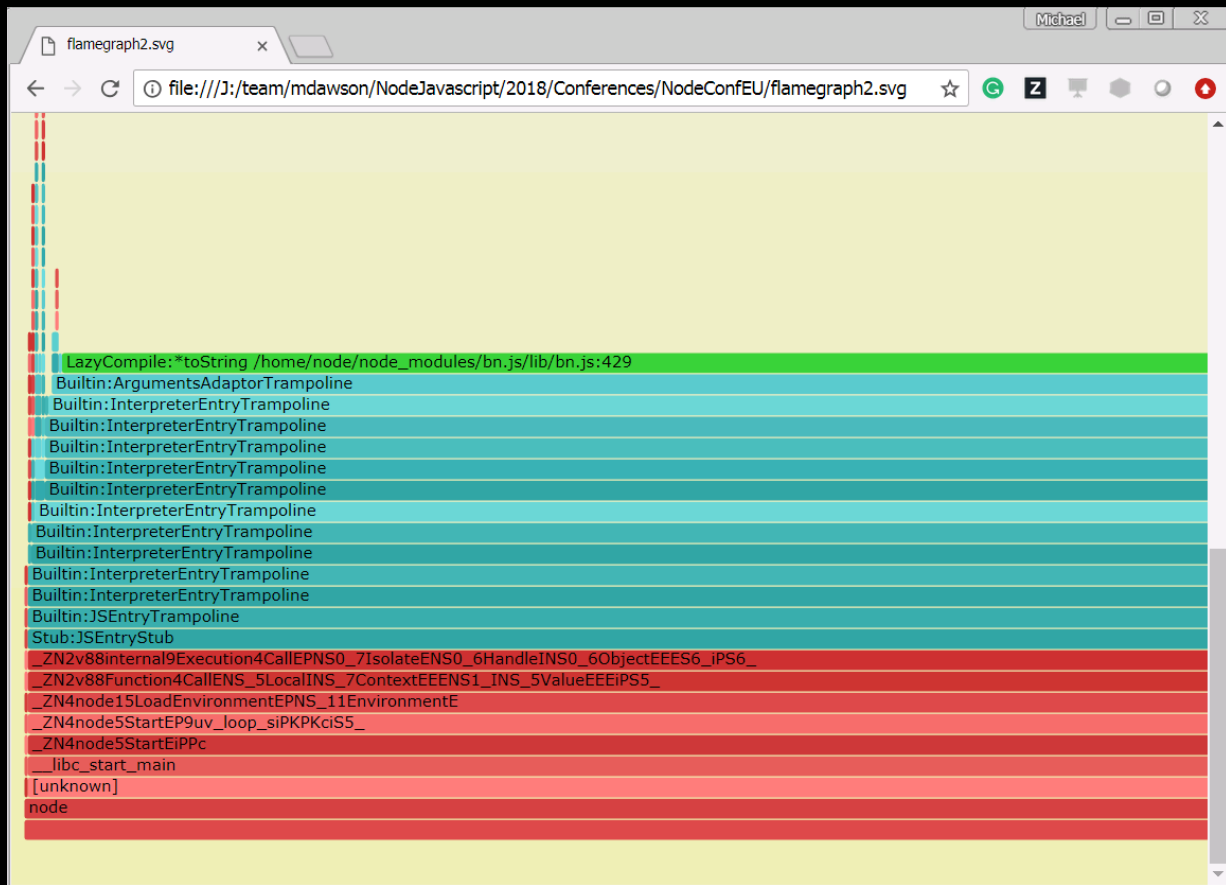
Scenario 3 – Flame Graphs



Scenario 3 – Flame Graphs



Scenario 3 – Flame Graphs



- 8.X missing info ?
- Importance of Tiers/testing

Recap

- Diagnostic Scenarios/Toolbox
- Diagnostics Working Group
- How this helps you

Help Out!

- Find an Initiative
- Make contact
- Contribute

Join us in creating the future of Node.js !

Copyrights and Trademarks

© IBM Corporation and Microsoft Corporation 2018. All Rights Reserved

IBM, the IBM logo, ibm.com are trademarks or registered trademarks of International Business Machines Corp., registered in many jurisdictions worldwide. Other product and service names might be trademarks of IBM or other companies.

A current list of IBM trademarks is available on the Web at
“Copyright and trademark information” at

www.ibm.com/legal/copytrade.shtml

Node.js is an official trademark of Joyent. IBM SDK for Node.js is not formally related to or endorsed by the official Joyent Node.js open source or commercial project.

Java, JavaScript and all Java-based trademarks and logos are trademarks or registered trademarks of Oracle and/or its affiliates.

Linux is a registered trademark of Linus Torvalds in the United States, other countries, or both.

Microsoft, Windows, Windows NT, and the Windows logo are trademarks of Microsoft Corporation in the United States, other countries, or both.

“TWITTER, TWEET, RETWEET and the Twitter logo are trademarks of Twitter, Inc. or its affiliates.”